

-iree-codegen-bufferize-copy-only-dispatches : Bufferize dispatches that copy to/from interfaces to convert to a `linalg.copy` op

-iree-codegen-canonicalize-scf-for : Adhoc canonicalization of selected loop-carried values/dependencies for `scf.for` ops

-iree-codegen-cleanup-buffer-alloc-view : Performs cleanups over HAL interface/buffer allocation/view operations

-iree-codegen-concretize-pad-result-shape : Concretizes `tensor.pad` op's result shape if its source op implements `OffsetSizeAndStrideOpInterface`.

-iree-codegen-convert-to-destination-passing-style : Transforms the code to make the dispatch use destination-passing style

Options

-use-war-for-cooperative-matrix-codegen : WAR for failure in Cooperative matrix codegen pipelines. See #10648.

-iree-codegen-decompose-linalg-generic : Decomposes `linalg` generic ops into individual ops

It is sometimes advantageous to operate on generic ops which contain at most one non-yield body operation. This is most often the case when needing to materialize individual ops (which some backends require). Note that this is often an extreme pessimization unless if part of a lowering flow which was designed for it.

Operates on tensor based `linalg` ops.

-iree-codegen-erase-hal-descriptor-type-from-memref : Erase `#hal.descriptor_type` from `MemRef` memory space

-iree-codegen-flatten-memref-subspan : Flatten n-D `MemRef` subspan ops to 1-D ones and fold byte offsets

-iree-codegen-fold-affinemem-in-distributed-loops : Fold `affine.mem` ops in distributed loops

-iree-codegen-fold-tensor-extract-op : Fold `tensor.extract` operations prior to lowering to LLVM

-iree-codegen-fuse-tensor-pad-with-consumer : Fuse `tensor.pad` op into its consumer op's tiled loop nest

-iree-codegen-gpu-tile-reduction : Pass to tile `linalg` reduction dimensions.

-iree-codegen-gpu-vectorization : Pass to convert `linalg` into `Vector`.

Options

-generate-contract : Try to convert reduction to `vector.contract`.

-max-vector-size : Max vector size allowed to avoid creating large vectors.

-iree-codegen-iree-comprehensive-bufferize : Convert from to Linalg ops on tensors to buffers

Options

-test-analysis-only : Only runs inplaceability analysis (for testing purposes only)
-print-conflicts : Annotates IR with RaW conflicts. Requires test-analysis-only.

-iree-codegen-memrefcopy-to-linalg : Convert memref.copy to linalg op

-iree-codegen-optimize-vector-transfer : Run optimization transformations on vector transfer operations

Options

-flatten : Flatten the vector type of vector transfers where possible (contiguous row-major data).

-iree-codegen-pad-dynamic-alloc : Pass to pad dynamic alloc into static one.

-iree-codegen-polynomial-approximation : Convert math operations to their polynomial approximation

-iree-codegen-reduction-to-gpu : Convert vector reduction to gpu ops.

-iree-codegen-rematerialize-parallel-ops : Pass to rematerialize and merge parallel ops to avoid creating temporary allocs.

-iree-codegen-remove-single-iteration-loop : Remove distributed loop with single iteration.

-iree-codegen-split-full-partial-transfer : Split a vector.transfer operation into an in-bounds (i.e., no out-of-bounds masking) fastpath and a slowpath.

Options

-split-transfers : Split vector transfers between slow (masked) and fast "(unmasked) variants. Possible options are:\n"\n\tnone [default]: keep unsplit vector.transfer and pay the price\n"\n\tlinalg-copy: use linalg.fill + linalg.generic for the slow path\n"\n\tvector-transfers: use extra small unmasked vector.transfers for"\n the slow path\n

-iree-codegen-test-partitionable-loops-interface : Test the PartitionableLoopsInterface

-iree-codegen-tile-and-distribute-to-workgroups : Tile and distribute operations to workgroups

-iree-codegen-type-propagation : Propagate the type of tensor to avoid load/stores of illegal bit widths

-iree-codegen-vectorize-tensor-pad : Vectorize a very specific form of tensor.pad with control flows

-iree-codegen-workgroup-specialization : Specialize workgroup distribution loops

-iree-convert-to-llvm : Perform final conversion from Linalg/HAL/Shape/Vector/Standard to LLVMIR dialect

Options

-reassociateFpReductions : Specifies if FP add and mult reductions can be reordered

-iree-convert-to-nvvm : Perform final conversion from builtin/GPU/HAL/standard dialect to LLVM and NVVM dialects

-iree-convert-to-rocdl : Perform final conversion from builtin/GPU/HAL/standard dialect to LLVM and ROC DL dialects

-iree-convert-to-spirv : Perform the final conversion to SPIR-V dialect

-iree-eliminate-empty-tensors : Eliminate tensor.empty ops to avoid buffer allocations

-iree-gpu-distribute-shared-memory-copy : Pass to distribute shared memory copies to threads.

-iree-gpu-multi-buffering : Pass to do multi buffering.

-iree-gpu-pipelining : Pass to do software pipelining.

Options

-epilogue-peeling : Try to use un-peeling epilogue when false, peeled epilogue o.w.

-iree-gpu-reduce-bank-conflicts : Pass to try to reduce the number of bank conflicts.

-iree-llvmcpu-assign-constant-ordinals : Assigns executable constant ordinals across all LLVMCPU variants.

-iree-llvmcpu-assign-import-ordinals : Assigns executable import ordinals across all LLVMCPU variants.

-iree-llvmcpu-check-ir-before-llvm-conversion : Checks CPU backend specific IR constraints (like no allocas)

-iree-llvmcpu-emit-vectorization-remarks : Emit vectorization remarks on Linalg ops

`-iree-llvmcpu-link-executables` : Links LLVMCPU HAL executables within the top-level program module.

`-iree-llvmcpu-lower-executable-target` : Lower executable target using an `IREE::HAL::DispatchLoweringPassPipeline`

`-iree-llvmcpu-materialize-encoding` : Materialize the encoding for tensor as specified by the backend

`-iree-llvmcpu-mmt4d-vector-lowering` : Apply vector lowering logic to vector ops

`-iree-llvmcpu-synchronize-symbol-visibility` : Synchronizes LLVM linkage with MLIR symbol visibility

`-iree-llvmcpu-unfuse-fma-pass` : Convert `llvm.fma` into unfused `mul` and `add` ops

`-iree-llvmcpu-vector-contract-custom-kernels` : Enable custom kernels (inline assembly or intrinsics) for some `vector.contract` ops

`-iree-llvmcpu-verify-linalg-transform-legality` : Verify that only supported IR constructs are passed to the compiler.

`-iree-llvmgpu-alloc` : Pass to create allocation for some values.

`-iree-llvmgpu-distribute` : Pass to distribute `foreachthread` ops.

`-iree-llvmgpu-lower-executable-target` : Perform lowering of executable target using one of the `IREE::HAL::DispatchLoweringPassPipeline`

`-iree-llvmgpu-tensor-pad` : Pass to pad out tensors up to static dimensions.

`-iree-llvmgpu-tensorcore-vectorization` : Pass to convert `linalg` into `Vector` and transform it to a form that can be lowered to GPU MMA ops

`-iree-llvmgpu-tile-and-distribute` : Pass to tile and distribute `linalg` ops within a workgroup.

`-iree-llvmgpu-tile-tensor` : Pass to tile `linalg` on tensor ops within a workgroup.

`-iree-llvmgpu-vector-lowering` : Pass to lower `Vector` ops before conversion to LLVM.

`-iree-llvmgpu-vector-to-gpu` : Pass to convert vector to gpu.

`-iree-spirv-annotate-winograd-loops` : Annotate innermost Winograd loops with `spirv distribute` attribute

`-iree-spirv-breakdown-large-vector` : Break down vectors not natively supported by SPIR-V

-iree-spirv-create-fast-slow-path : Create separate fast and slow paths to handle padding

-iree-spirv-distribute : Distribute tiled loop nests to invocations

-iree-spirv-emulate-i64 : Emulate 64-bit integer ops with 32-bit integer ops

-iree-spirv-lower-executable-target-pass : Lower the executable target to SPIR-V using one of the `IREE::HAL::DispatchLoweringPassPipeline`

-iree-spirv-map-memref-storage-class : Map MemRef memory spaces to SPIR-V storage classes

-iree-spirv-tile : Tile Linalg ops with tensor semantics to invocations

-iree-spirv-tile-and-distribute : Tile and distribute Linalg ops with buffer semantics to invocations

-iree-spirv-tile-and-promote : Promote tiled Linalg ops with buffer semantics to use workgroup memory and then tile to invocations

Options

-promote-c : Promote C matrix to use shared memory
-skip-thread : Skip tiling and distributing to GPU threads

-iree-spirv-tile-to-cooperative-ops : Tile Linalg ops with buffer semantics to subgroups and vectorize to vector ops suitable for lowering to SPIR-V cooperative ops

-iree-spirv-vector-to-gpu-subgroup-mma-ops : Pass to convert vector ops to GPU subgroup MMA ops.

-iree-spirv-vectorize : Vectorize Linalg ops with buffer semantics

-iree-spirv-vectorize-load-store : Vectorize load/store of memrefs for better memory access

-iree-spirv-vectorize-to-cooperative-ops : Tile Linalg ops with buffer semantics to subgroups and vectorize to vector ops suitable for lowering to SPIR-V cooperative ops

-iree-test-llvmgpu-legalize-ops : Test pass for several legalization patterns.

-iree-transform-dialect-interpreter : Pass to apply transform dialect operations.

Options

-transform-file-name : Optional filename containing a transform dialect specification to apply. If left empty, the IR is assumed to contain one top-level

transform dialect operation somewhere in the module.

-debug-payload-root-tag : Select the operation with 'transform.iree_tag' attribute having the given value as payload IR root. This allows user control on what operation to transform in debug mode, without requiring intimate knowledge of the IREE nested pass pipeline.\nIf empty (normal operation mode), select the pass anchor operation in the IREE pipeline, as the payload IR root.

-debug-transform-root-tag : Select the operation with 'transform.iree_tag' attribute having the given value as container IR for top-level transform ops. This allows user control on what transformation to apply in debug mode, without requiring intimate knowledge of the IREE nested pass pipeline.\nIf empty (normal operation mode), select the container of the top-level transform op.

-iree-vmvx-assign-constant-ordinals : Assigns executable constant ordinals across all VMVX variants.

-iree-vmvx-link-executables : Links VMVX HAL executables within the top-level program module.

-iree-vmvx-lower-linalg-microkernels : Lowers linalg ops to the VMVX microkernel library

Options

-warn-on-unconverted : Warns on any unconverted linalg ops which remain live

-iree-vmvx-materialize-encoding : Materialize the encoding for tensor as specified by the backend

-iree-wgsl-replace-push-constants : Replaces push constant loads with binding loads for when using WGSL without push constant support

-iree-workgroup-swizzle : swizzle the workgroup ids for better cache reuse

Options

-logTile : pass the tile value for unit testing